


Data Mining (Minería de Datos)

The k-NN technique




Rodrigo Manzananas

Grupo de Meteorología
Univ. de Cantabria – CSIC
MACC / IFCA

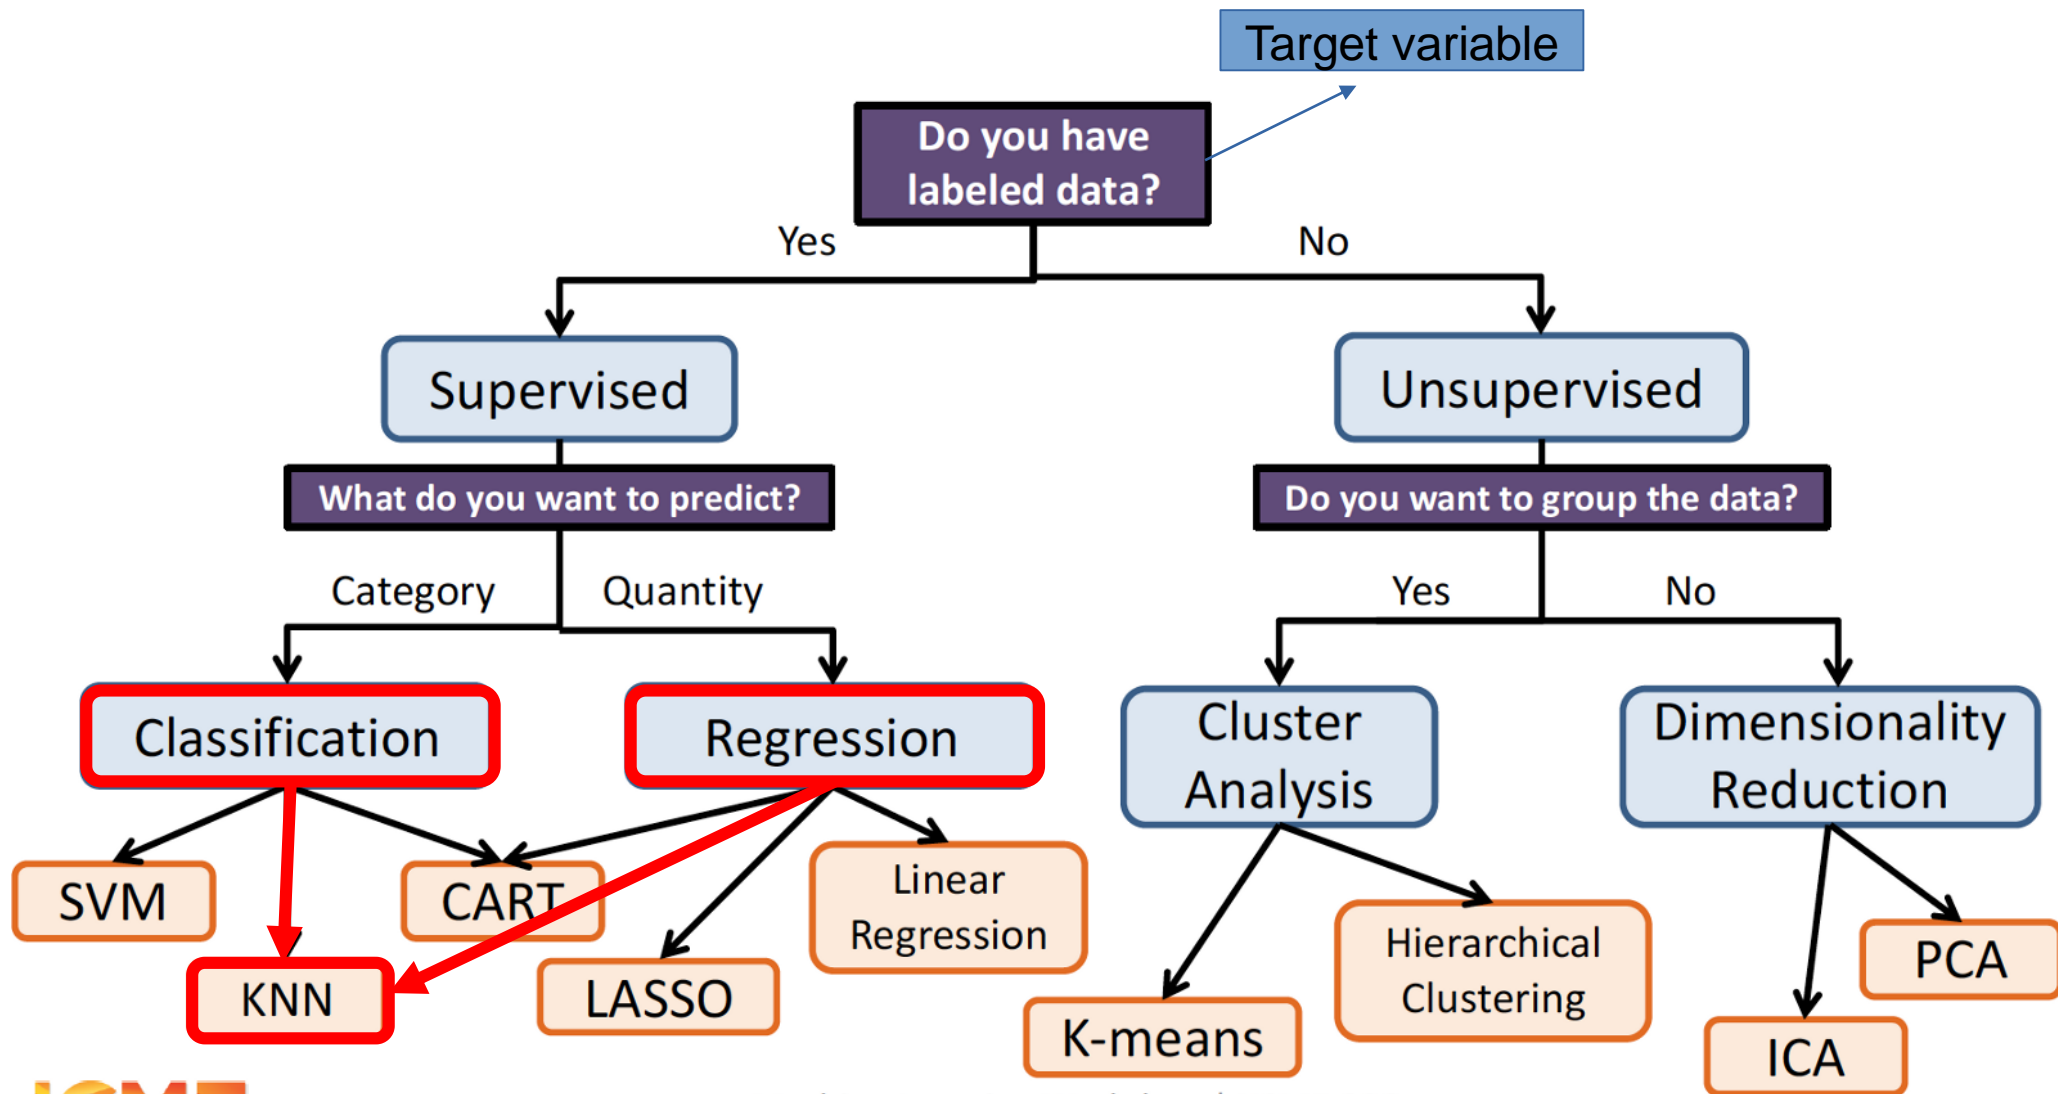


Types of Machine Learning



NOTA: Las líneas de código de R en esta presentación se muestran sobre un fondo gris

Oct	30	Aplazada (sesión de refuerzo)
Nov	6	Presentación, introducción y perspectiva histórica
	8	Paradigmas, problemas canónicos y data challenges
	13	Reglas de asociación
	15	Práctica: Reglas de asociación
	20	Evaluación, sobreajuste y cross-validación
	22	Práctica: Cross-validación
	27	Árboles de clasificación
	29	Práctica: Árboles de clasificación
		T01. Datos discretos
Dic	4	Técnicas de vecinos cercano (k-NN)
	11	Práctica: Vecinos cercanos
	13	Reducción de dimensión lineal
	18	Práctica: LDA y PCA
	20	Reducción no lineal
		T02. Clasificación
Ene	8	Árboles de clasificación y regresión (CART)
	10	Práctica: CART
	15	Ensembles: Bagging and Boosting
	17	Práctica: Random forests
		T03. Predicción
	22	Práctica: Gradient boosting
	24a	Técnicas de agrupamiento
	24b	Práctica: Técnicas de agrupamiento
	29a	Práctica: El paquete CARET
	29b	Examen



How does it work?

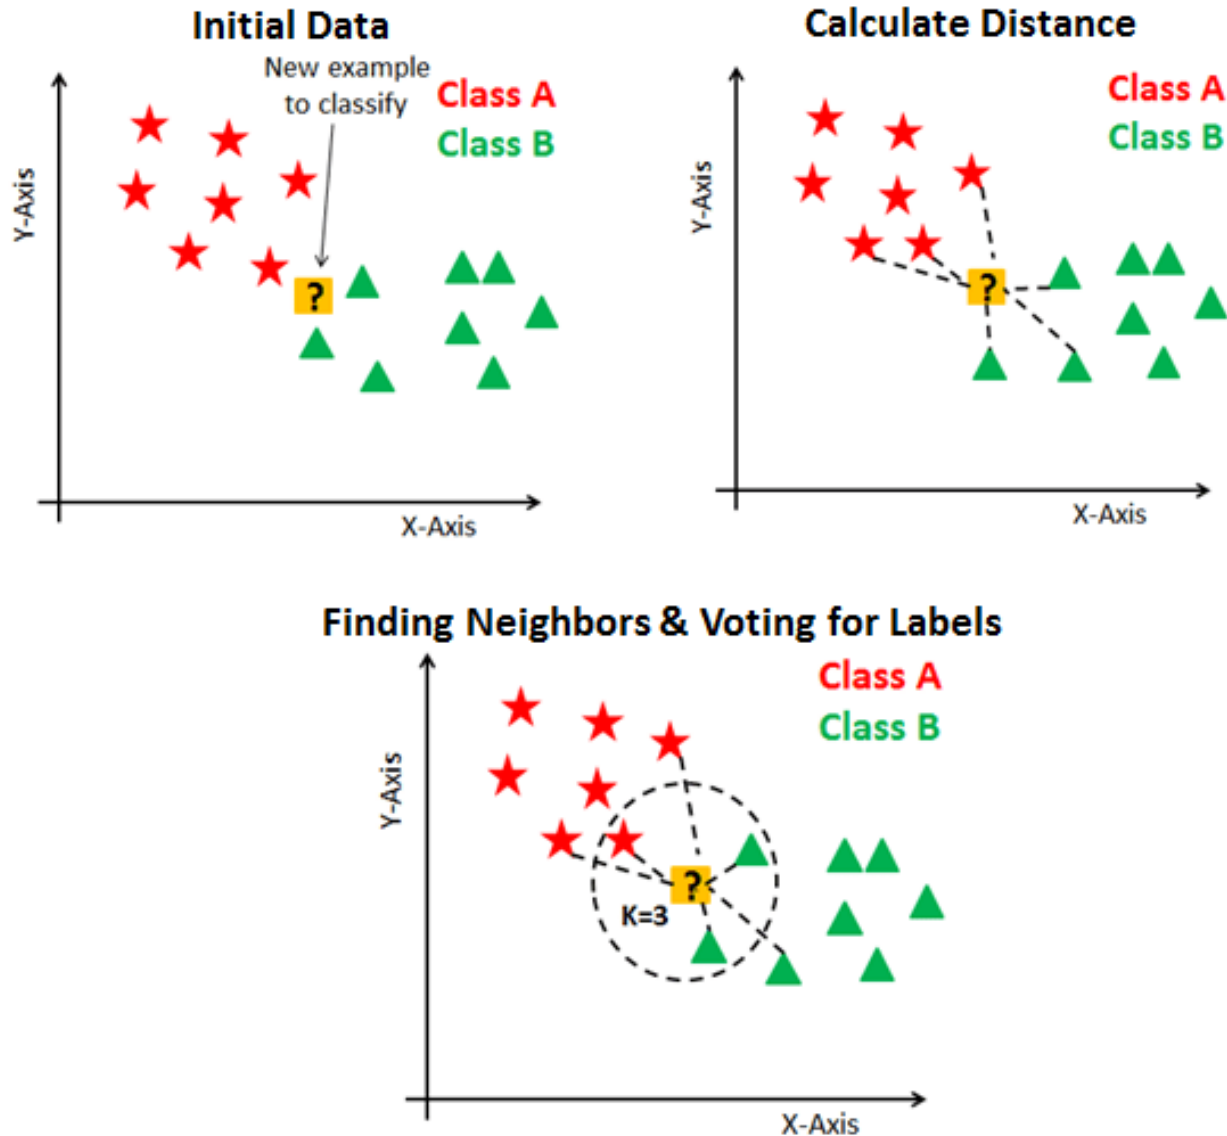
“... dime con quién vas y te diré quién eres ...”

Non-parametric:

No assumption is made on the underlying data distribution

Lazy (or instance-based) learning:

There is no explicit training phase. All the training data is needed during the testing phase



How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

1) Computing distances

$$d_{151,1} = \sqrt{(5.4 - 5.1)^2 + (2.7 - 3.5)^2 + (4.6 - 1.4)^2 + (1.4 - 0.2)^2} = 3.52$$

How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

1) Computing distances

$$d_{151,1} = \sqrt{(5.4 - 5.1)^2 + (2.7 - 3.5)^2 + (4.6 - 1.4)^2 + (1.4 - 0.2)^2} = 3.52$$

$$d_{151,2} = 3.47$$

$$d_{151,3} = 3.62$$

$$d_{151,55} = 1.11$$

$$d_{151,56} = 0.35$$

$$d_{151,57} = 1.10$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,150} = 0.87$$

How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

1) Computing distances

$$d_{151,1} = \sqrt{(5.4 - 5.1)^2 + (2.7 - 3.5)^2 + (4.6 - 1.4)^2 + (1.4 - 0.2)^2} = 3.52$$

$$d_{151,2} = 3.47$$

$$d_{151,3} = 3.62$$

$$d_{151,55} = 1.11$$

$$d_{151,56} = 0.35$$

$$d_{151,57} = 1.10$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,150} = 0.87$$

2) Ordering distances

$$d_{151,56} = 0.35$$

$$d_{151,150} = 0.87$$

$$d_{151,57} = 1.10$$

$$d_{151,55} = 1.11$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,2} = 3.47$$

$$d_{151,1} = 3.52$$

$$d_{151,3} = 3.62$$

How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

1) Computing distances

$$d_{151,1} = \sqrt{(5.4 - 5.1)^2 + (2.7 - 3.5)^2 + (4.6 - 1.4)^2 + (1.4 - 0.2)^2} = 3.52$$

$$d_{151,2} = 3.47$$

$$d_{151,3} = 3.62$$

$$d_{151,55} = 1.11$$

$$d_{151,56} = 0.35$$

$$d_{151,57} = 1.10$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,150} = 0.87$$

2) Ordering distances 3) Classification based on NN

$d_{151,56} = 0.35$

 $\xrightarrow{k=1}$
 $151 = \text{Versicolor}$

$$d_{151,150} = 0.87$$

$$d_{151,57} = 1.10$$

$$d_{151,55} = 1.11$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,2} = 3.47$$

$$d_{151,1} = 3.52$$

$$d_{151,3} = 3.62$$

How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

1) Computing distances

$$d_{151,1} = \sqrt{(5.4 - 5.1)^2 + (2.7 - 3.5)^2 + (4.6 - 1.4)^2 + (1.4 - 0.2)^2} = 3.52$$

$$d_{151,2} = 3.47$$

$$d_{151,3} = 3.62$$

$$d_{151,55} = 1.11$$

$$d_{151,56} = 0.35$$

$$d_{151,57} = 1.10$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,150} = 0.87$$

2) Ordering distances

$d_{151,56} = 0.35$
 $d_{151,150} = 0.87$

$d_{151,57} = 1.10$
 $d_{151,55} = 1.11$
 $d_{151,148} = 1.42$
 $d_{151,149} = 1.61$
 $d_{151,2} = 3.47$
 $d_{151,1} = 3.52$
 $d_{151,3} = 3.62$

3) Classification based on NN

$k = 1 \rightarrow 151 = \text{Versicolor}$

$k = 2 \rightarrow 151 = ?$

How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

1) Computing distances

$$d_{151,1} = \sqrt{(5.4 - 5.1)^2 + (2.7 - 3.5)^2 + (4.6 - 1.4)^2 + (1.4 - 0.2)^2} = 3.52$$

$$d_{151,2} = 3.47$$

$$d_{151,3} = 3.62$$

$$d_{151,55} = 1.11$$

$$d_{151,56} = 0.35$$


$$d_{151,57} = 1.10$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,150} = 0.87$$

2) Ordering distances



$$d_{151,55} = 1.11$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,2} = 3.47$$

$$d_{151,1} = 3.52$$

$$d_{151,3} = 3.62$$

3) Classification based on NN

How does it work?

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
...					
55	6.5	2.8	4.6	1.5	versicolor
56	5.7	2.8	4.5	1.3	versicolor
57	6.3	3.3	4.7	1.6	versicolor
...					
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica
...					
151	5.4	2.7	4.6	1.4	?

1) Computing distances

$$d_{151,1} = \sqrt{(5.4 - 5.1)^2 + (2.7 - 3.5)^2 + (4.6 - 1.4)^2 + (1.4 - 0.2)^2} = 3.52$$

$$d_{151,2} = 3.47$$

$$d_{151,3} = 3.62$$

$$d_{151,55} = 1.11$$

$$d_{151,56} = 0.35$$

$$d_{151,57} = 1.10$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,150} = 0.87$$

2) Ordering distances

$d_{151,56} = 0.35$
$d_{151,150} = 0.87$
$d_{151,57} = 1.10$

$$d_{151,55} = 1.11$$

$$d_{151,148} = 1.42$$

$$d_{151,149} = 1.61$$

$$d_{151,2} = 3.47$$

$$d_{151,1} = 3.52$$

$$d_{151,3} = 3.62$$

3) Classification based on NN

$k = 1$	\rightarrow	151 = Versicolor
$k = 2$	\rightarrow	151 = ?
$k = 3$	\rightarrow	151 = Versicolor

Pros:

- Easy to understand
- Versatile: Classification and regression problems
- High accuracy (benchmark method)

Cons:

- High memory requirements, computationally expensive
- Sensitive to scale of the data
- Can suffer from biases towards skewed distributions
- Performance can be severely degraded in high-dimensional problems

Applications:

- Economic sciences: concession of loans
- Political sciences: classifying potential voters
- Handwriting detection (e.g. OCR)
- Image/video recognition
- Genetics

Fitting the method

Distance metric

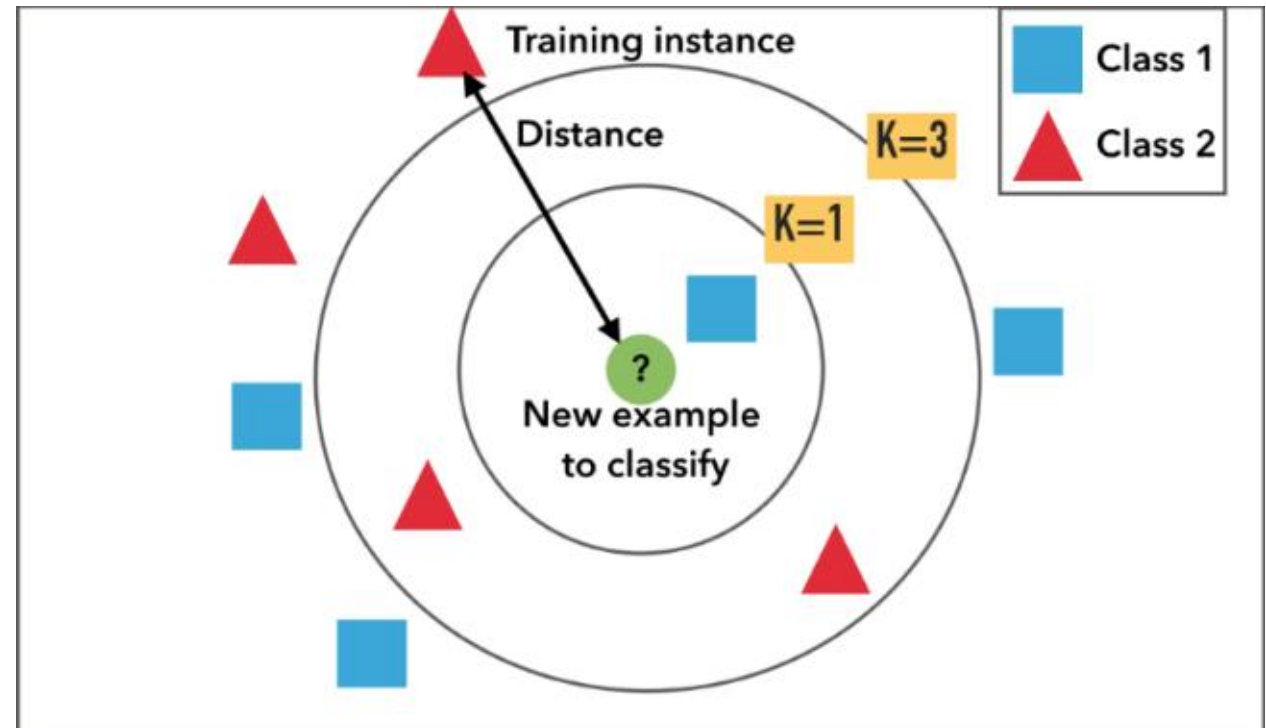
Different distances are used, depending on the application. *Euclidean* is the most common

Number of neighbors (k)

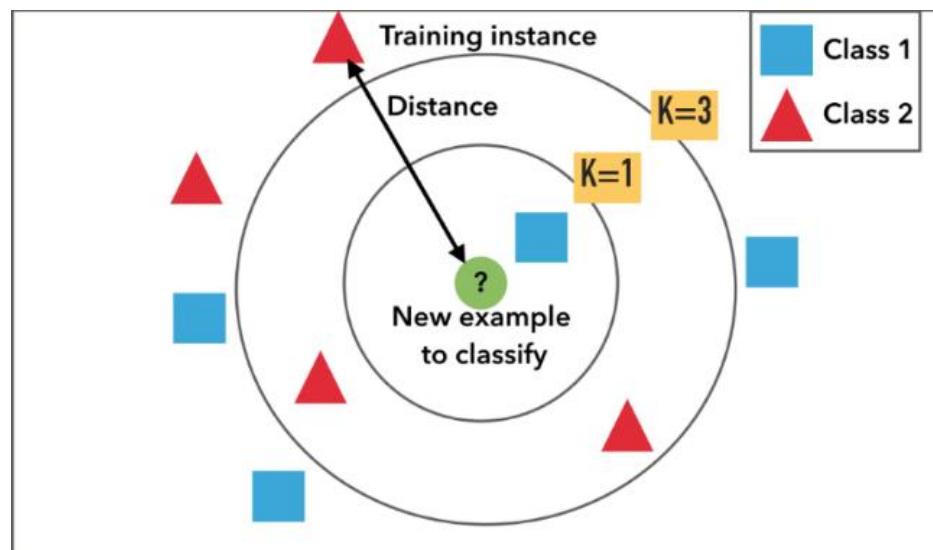
This is the unique model parameter. Must be properly chosen

Classifying criterion

- Majority vote
- Weighted vote
- Random
- etc.



Distance metric



Minkowsky:

$$D(x, y) = \left(\sum_{i=1}^m |x_i - y_i|^r \right)^{1/r}$$

Euclidean:

$$D(x, y) = \sqrt{\sum_{i=1}^m (x_i - y_i)^2}$$

Manhattan / city-block:

$$D(x, y) = \sum_{i=1}^m |x_i - y_i|$$

Camberra:


$$D(x, y) = \sum_{i=1}^m \frac{|x_i - y_i|}{|x_i + y_i|}$$

Chebychev:

$$D(x, y) = \max_{i=1}^m |x_i - y_i|$$

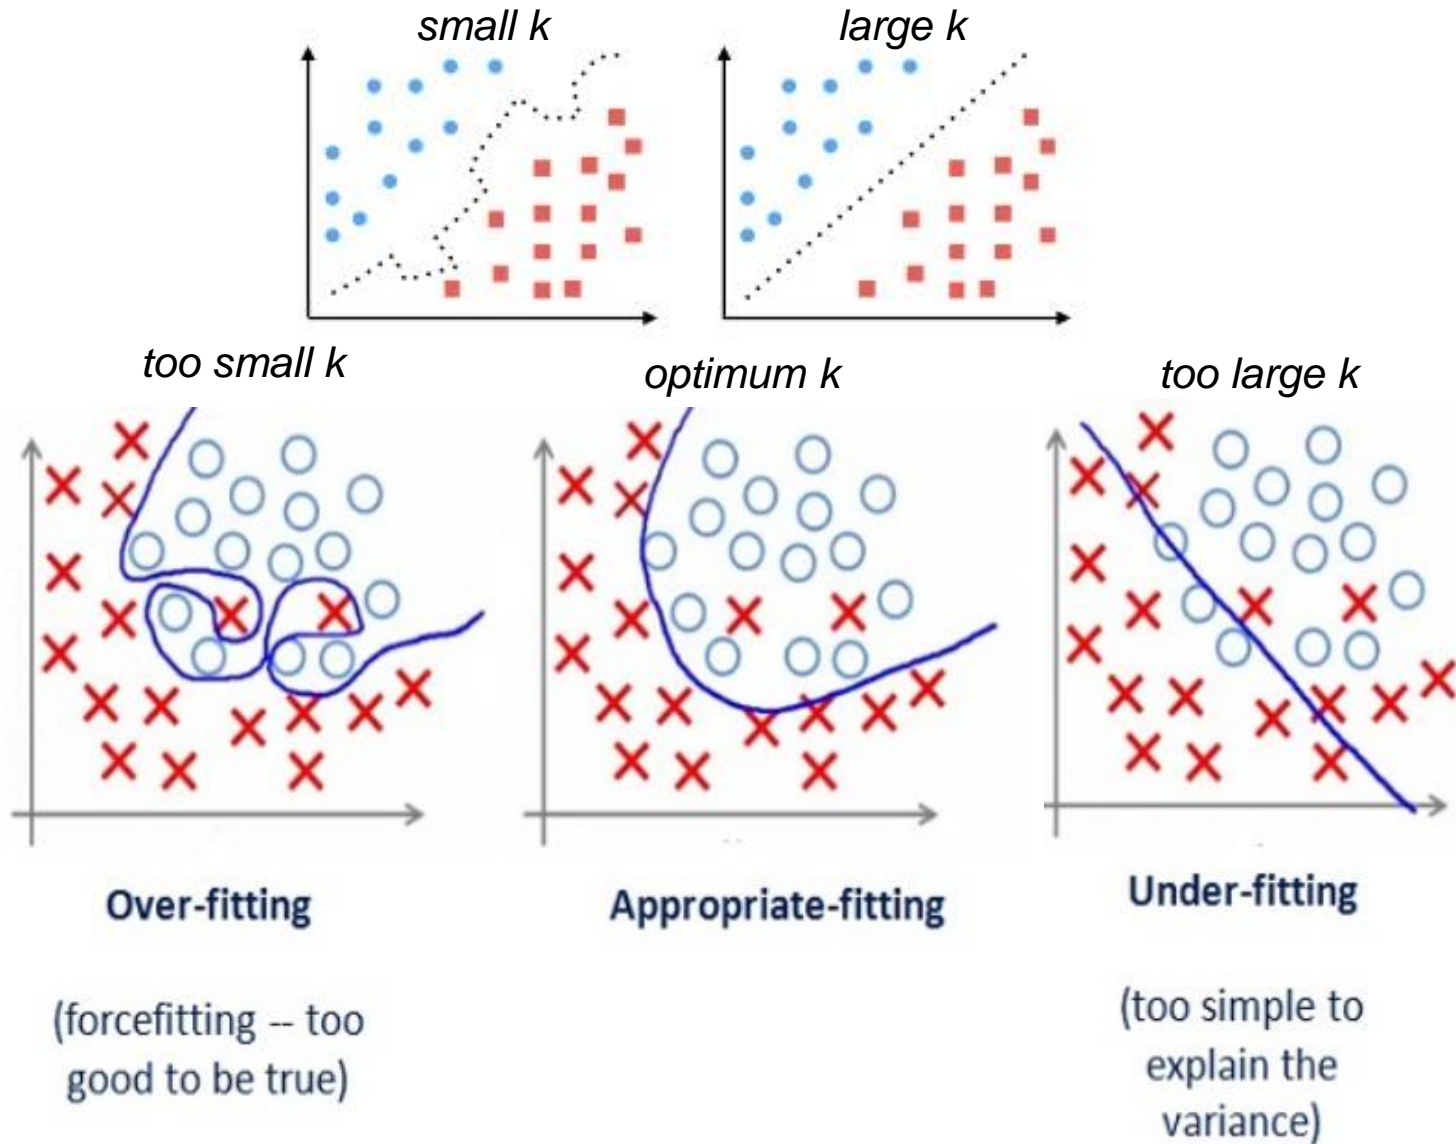
Number of neighbors (k)

Which one classifies better?




Number of neighbors (k)

Which one classifies better?



Number of neighbors (k)

Cross-validation is needed to find the optimal k



Examples in R

Based on the iris dataset, classify the following new instance:

(sepal l., sepal w., petal l., petal w.) = (5.4, 2.7, 4.6, 1.4)

```
# new instance  
d.new = c(5.4, 2.7, 4.6, 1.4)
```

Examples in R

Based on the iris dataset, classify the following new instance:

(sepal l., sepal w., petal l., petal w.) = (5.4, 2.7, 4.6, 1.4)

```
# new instance  
d.new = c(5.4, 2.7, 4.6, 1.4)
```

```
# euclidean distance between the new instance and all the others  
eucli = c()  
for (i in 1:nrow(iris)) {  
  eucli[i] = sqrt(sum((d.new - iris[i,-5])^2))  
}
```

Examples in R

Based on the iris dataset, classify the following new instance:

(sepal l., sepal w., petal l., petal w.) = (5.4, 2.7, 4.6, 1.4)

```
# new instance  
d.new = c(5.4, 2.7, 4.6, 1.4)
```

```
# euclidean distance between the new instance and all the others  
eucli = c()  
for (i in 1:nrow(iris)) {  
  eucli[i] = sqrt(sum((d.new - iris[i,-5])^2))  
}
```

```
## ordering distances  
ind.sort = sort(eucli, index.return = T)
```

Examples in R

Based on the iris dataset, classify the following new instance:

(sepal l., sepal w., petal l., petal w.) = (5.4, 2.7, 4.6, 1.4)

```
# new instance  
d.new = c(5.4, 2.7, 4.6, 1.4)
```

```
# euclidean distance between the new instance and all the others  
eucli = c()  
for (i in 1:nrow(iris)) {  
  eucli[i] = sqrt(sum((d.new - iris[i,-5])^2))  
}
```

```
## ordering distances  
ind.sort = sort(eucli, index.return = T)
```

```
# classifying based on the nearest neighbor  
pred.k1 = iris$Species[ind.sort$ix[1]]
```

Examples in R

Based on the iris dataset, classify the following new instance:

(sepal l., sepal w., petal l., petal w.) = (5.4, 2.7, 4.6, 1.4)

```
# new instance  
d.new = c(5.4, 2.7, 4.6, 1.4)
```

```
# euclidean distance between the new instance and all the others  
eucli = c()  
for (i in 1:nrow(iris)) {  
  eucli[i] = sqrt(sum((d.new - iris[i,-5])^2))  
}
```

```
## ordering distances  
ind.sort = sort(eucli, index.return = T)
```

```
# classifying based on the nearest neighbor  
pred.k1 = iris$Species[ind.sort$ix[1]]
```

```
# classifying based on the 10 nearest neighbors  
pred.k10 = iris$Species[ind.sort$ix[1:10]]  
summary(pred.k10)
```


Examples in R

Divide iris into train y test (75% and 25% of the total dataset, respectively) and find the test error for the nearest neighbor method. Use the function “knn” from package “class”

Examples in R

Divide iris into train y test (75% and 25% of the total dataset, respectively) and find the test error for the nearest neighbor method. Use the function “knn” from package “class”

```
# train/test division  
n = nrow(iris)  
indtrain = sample(1:n, round(0.75*n))  
indtest = setdiff(1:n, indtrain)  
iris.train = iris[indtrain,]  
iris.test = iris[indtest,]
```

Examples in R

Divide iris into train y test (75% and 25% of the total dataset, respectively) and find the test error for the nearest neighbor method. Use the function “knn” from package “class”

```
# train/test division
n = nrow(iris)
indtrain = sample(1:n, round(0.75*n))
indtest = setdiff(1:n, indtrain)
iris.train = iris[indtrain,]
iris.test = iris[indtest,]
```

```
# classifying using the nearest neighbor method
library(class)
pred = knn(train = iris.train[,-5], test = iris.test[,-5], cl = iris.train$Species, k = 1)
```

Examples in R

Divide iris into train y test (75% and 25% of the total dataset, respectively) and find the test error for the nearest neighbor method. Use the function “knn” from package “class”

```
# train/test division
n = nrow(iris)
indtrain = sample(1:n, round(0.75*n))
indtest = setdiff(1:n, indtrain)
iris.train = iris[indtrain,]
iris.test = iris[indtest,]
```

```
# classifying using the nearest neighbor method
library(class)
pred = knn(train = iris.train[,-5], test = iris.test[,-5], cl = iris.train$Species, k = 1)
```

```
# validating method
table(pred, iris.test$Species)
pred      setosa versicolor virginica
setosa    11      0      0
versicolor 0     14      1
virginica  0      1     11
acc.class(pred, iris.test$Species)
```

```
# evaluation function
acc.class = function(x, y) {
  stopifnot(length(x) == length(y))
  return(sum(diag(table(x, y))) / length(x))
}
```

Examples in R

Use the package “caret” (method “knn”) to find the optimal k . To do so, check how the test error varies with increasing k (for values from 1 to 50) under a hold-out cross-validation scheme.


```
library(caret)
# defining hold-out cross-validation
trctrl = trainControl(method = "cv", number = 2)
# searching the optimal  $k$ 
knn.fit = train(Species ~ ., iris.train,
                method = "knn",
                trControl = trctrl,
                tuneGrid = expand.grid(k = 1:50))
plot(knn.fit)
```

)

Examples in R

Use the package “caret” (method “knn”) to find the optimal k . To do so, check how the test error varies with increasing k (for values from 1 to 50) under a hold-out cross-validation scheme.

```
library(caret)
# defining hold-out cross-validation
trctrl = trainControl(method = "cv", number = 2)
# searching the optimal  $k$ 
knn.fit = train(Species ~ ., iris.train,
                method = "knn",
                trControl = trctrl,
                tuneGrid = expand.grid(k = 1:50))
plot(knn.fit)
```




Examples in R

Use the package “caret” (method “knn”) to find the optimal k . To do so, check how the test error varies with increasing k (for values from 1 to 50) under a hold-out cross-validation scheme.

```
library(caret)
# defining hold-out cross-validation
trctrl = trainControl(method = "cv", number = 2)
# searching the optimal  $k$ 
knn.fit = train(Species ~ ., iris.train,
  method = "knn",
  trControl = trctrl,
  tuneGrid = expand.grid(k = 1:50))
plot(knn.fit)
```

```
# predicting in test with the optimal  $k$ 
pred = predict(knn.fit, iris.test)
acc = acc.class(pred, iris.test$Species)
```




k-NN for regression

Aim:
Predicting a continuous target variable

What do we need?
An inference criterion: It can be a simple mean, a particular percentile, etc.

To take into account:
Predictor variables covering larger ranges may have more weight in the search of neighbours. Rescaling the predictor data is recommended to make the distance metric more meaningful

$$Z = \frac{X - \mu}{\sigma}$$



	Ozone	Solar.R	Wind	Temp
1	41	190	7.4	67
2	36	118	8.0	72
3	12	149	12.6	74
4	18	313	11.5	62
5	25	297	14.3	56
...				
500	23	234	9.3	65
501	45	321	16.7	?

Examples in R

For regression, we will work with the dataset “carseats” (included in the package “ISLR”). Our target variable will be “Sales”. First, we will remove all the categorical variables from the dataset, retaining only the continuous ones. We will use the function “knn.reg” from the package “FNN”. As you did for the case of classification, divide the total dataset in 75% for train and 25% for test and see how the test error varies with k

Examples in R

For regression, we will work with the dataset “carseats” (included in the package “ISLR”). Our target variable will be “Sales”. First, we will remove all the categorical variables from the dataset, retaining only the continuous ones. We will use the function “knn.reg” from the package “FNN”. As you did for the case of classification, divide the total dataset in 75% for train and 25% for test and see how the test error varies with k

```
library(ISLR)
attach(Carseats)
dataset = Carseats[, -c(7,10,11)]

# evaluation function
rmse <- function(x, y) {
  sqrt(mean((x - y)^2))
}

# train/test division
n = nrow(dataset)
indtrain = sample(1:n, round(0.75*n));
dataset.train = dataset[indtrain, ]
indtest = setdiff(1:n, indtrain);
dataset.test = dataset[indtest, ]
```

Examples in R

For regression, we will work with the dataset “carseats” (included in the package “ISLR”). Our target variable will be “Sales”. First, we will remove all the categorical variables from the dataset, retaining only the continuous ones. We will use the function “knn.reg” from the package “FNN”. As you did for the case of classification, divide the total dataset in 75% for train and 25% for test and see how the test error varies with k

```
library(ISLR)
attach(Carseats)
dataset = Carseats[, -c(7,10,11)]

# evaluation function
rmse <- function(x, y) {
  sqrt(mean((x - y)^2))
}

# train/test division
n = nrow(dataset)
indtrain = sample(1:n, round(0.75*n));
dataset.train = dataset[indtrain, ]
indtest = setdiff(1:n, indtrain);
dataset.test = dataset[indtest, ]
```

```
# test error as a function of k
library(FNN)
kmax = 50
test.err = c()
for (k in 1:kmax) {
  pred = knn.reg(dataset.train[, -1], dataset.test[, -1],
    dataset.train$Sales, k = k)
  test.err[k] = rmse(pred$pred, as.numeric(dataset.test$Sales))
}
plot(1:kmax, test.err, type = "o", pch = 19,
  xlab = "k", ylab = "RMSE"); grid()
```

Examples in R


For regression, we will work with the dataset “carseats” (included in the package “ISLR”). Our target variable will be “Sales”. First, we will remove all the categorical variables from the dataset, retaining only the continuous ones. We will use the function “knn.reg” from the package “FNN”. As you did for the case of classification, divide the total dataset in 75% for train and 25% for test and see how the test error varies with k

```
library(ISLR)
attach(Carseats)
dataset = Carseats[, -c(7,10,11)]

# evaluation function
rmse <- function(x, y) {
  sqrt(mean((x - y)^2))
}

# train/test division
n = nrow(dataset)
indtrain = sample(1:n, round(0.75*n));
dataset.train = dataset[indtrain, ]
indtest = setdiff(1:n, indtrain);
dataset.test = dataset[indtest, ]
```

```
# test error as a function of k
library(FNN)
kmax = 50
test.err = c()
for (k in 1:kmax) {
  pred = knn.reg(dataset.train[, -1], dataset.test[, -1],
    dataset.train$Sales, k = k)
  test.err[k] = rmse(pred$pred, as.numeric(dataset.test$Sales))
}
plot(1:kmax, test.err, type = "o", pch = 19,
  xlab = "k", ylab = "RMSE"); grid()
```



Examples in R

Let's continue by seeing the effect of properly rescaling the predictor data. Use the function "scale"

```
# predictor ranges  
boxplot(dataset[, -1])
```




Examples in R

Let's continue by seeing the effect of properly rescaling the predictor data. Use the function "scale"

```
# predictor ranges  
boxplot(dataset[, -1])
```

```
# test error as a function of k (for standardized data)  
test.err2 = c()  
for (k in 1:kmax) {  
  pred = knn.reg(scale(dataset.train[, -1]),  
    scale(dataset.test[, -1]), dataset.train$Sales, k = k)  
  test.err2[k] = rmse(pred$pred,  
    as.numeric(dataset.test$Sales))  
}
```




Examples in R

Let's continue by seeing the effect of properly rescaling the predictor data. Use the function "scale"

```
# predictor ranges  
boxplot(dataset[, -1])
```

```
# test error as a function of k (for standardized data)  
test.err2 = c()  
for (k in 1:kmax) {  
  pred = knn.reg(scale(dataset.train[, -1]),  
    scale(dataset.test[, -1]), dataset.train$Sales, k = k)  
  test.err2[k] = rmse(pred$pred,  
    as.numeric(dataset.test$Sales))  
}
```

```
# plotting results  
matplot(1:kmax, cbind(test.err, test.err2),  
  type = "o", pch = 19, lty = 1,  
  col = c("black", "red"), xlab = "k", ylab = "RMSE")  
legend("topright", c("without scaling", "with scaling"),  
  lty = 1, col = c("black", "red"))  
grid()
```



Examples in R


Do the same exercise, but this time using “caret”. Recall to standardize your predictor data to obtain meaningful results.

Examples in R

Do the same exercise, but this time using “caret”. Recall to standardize your predictor data to obtain meaningful results

```
# defining hold-out cross-validation
trctrl <- trainControl(method = "cv", number = 2)
# searching the optimal k (with standardized data)
knn.fit <- train(Sales ~ ., data = dataset.train,
  method = "knn",
  trControl = trctrl,
  preProcess = c("center", "scale"),
  tuneGrid = expand.grid(k = 1:50))
plot(knn.fit)
```

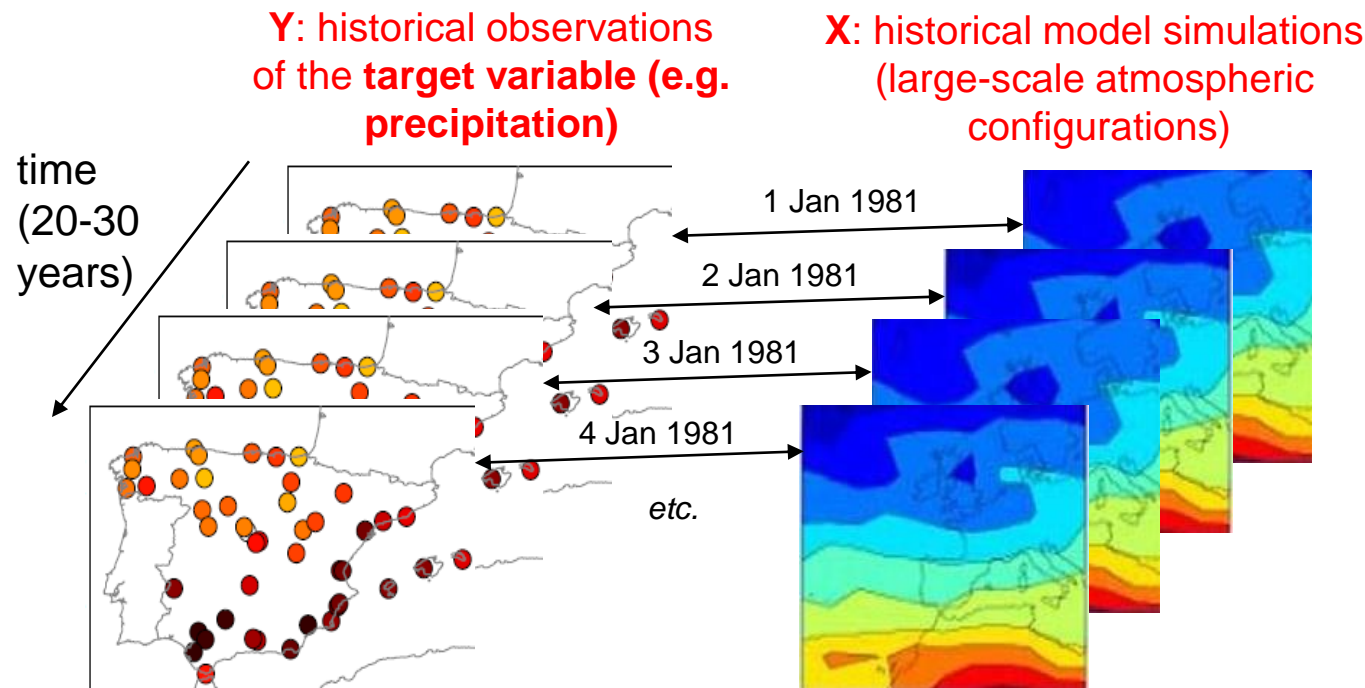
```
# predicting in test with the optimal k
pred = predict(knn.fit, dataset.test)
rmse(pred, dataset.test$Sales)
```




k-NN in meteorology

The analog technique (Lorenz, 1969): Similar atmospheric patterns lead to similar meteorological conditions


**Problem: Y' (prediction)
for 26 Mar 2046?**



1) Take X' for 26 Mar 2046: X_{2046}




2) Search the nearest neighbor/s to X_{2046} within X



X (3-Jan-1981)

3) Infer a prediction based on the observed values in the days selected in 2)




In meteorology, two important factors must be taken into account for the application of k-NN technique:

- 1) Predictor scaling
- 2) High dimensionality (the curse of dimensionality)

The curse of dimensionality

(David Scott, *Multivariate Density Estimation*, Wiley, 1992)




Volume of sphere relative to cube in d dimensions?

Dimension	2
Rel. vol.	0.79

The curse of dimensionality

(David Scott, *Multivariate Density Estimation*, Wiley, 1992)




Volume of sphere relative to cube in d dimensions?

Dimension	2	3
Rel. vol.	0.79	0.53

The curse of dimensionality

(David Scott, *Multivariate Density Estimation*, Wiley, 1992)




Volume of sphere relative to cube in d dimensions?

Dimension	2	3	4
Rel. vol.	0.79	0.53	0.31

The curse of dimensionality

(David Scott, *Multivariate Density Estimation*, Wiley, 1992)




Volume of sphere relative to cube in d dimensions?

Dimension	2	3	4	5
Rel. vol.	0.79	0.53	0.31	0.16

The curse of dimensionality

(David Scott, *Multivariate Density Estimation*, Wiley, 1992)




Volume of sphere relative to cube in d dimensions?

Dimension	2	3	4	5	6
Rel. vol.	0.79	0.53	0.31	0.16	0.08

The curse of dimensionality


(David Scott, *Multivariate Density Estimation*, Wiley, 1992)



Volume of sphere relative to cube in d dimensions?



Dimension	2	3	4	5	6	7
Rel. vol.	0.79	0.53	0.31	0.16	0.08	0.04

As the dimensionality increases, a larger percentage of the training data resides in the corners of the feature space. Therefore, k-NN is unhelpful in high dimensional problems because there is little difference between the nearest and the farthest neighbor



The curse of dimensionality

10% Neighborhood



The amount of training data needed to cover 10% of the feature range grows exponentially with the number of dimensions

Dimensionality reduction techniques (e.g. PCA ~ effective degrees of freedom) should be applied prior to using k-NN in order to help make the distance metric more meaningful