

# Semántica, Datos Conectados y Minería de Datos Textual

## Minería de textos y minería Web

Cristina Tîrnăucă

Dept. Matesco, Universidad de Cantabria

Facultad de Ciencias – Máster en Data Science

## Text Mining

### Definition

Text mining is the process of **deriving high-quality information** from **text**.

### Text analysis involves

- information retrieval (to find relevant documents),
- lexical analysis (to study word frequency distributions),
- syntactic parsing
- pattern recognition,
- tagging/annotation,
- information extraction,
- data mining techniques (including link and association analysis, visualization, and predictive analytics).

## Web Mining

### Definition

Web mining is the process of **finding patterns** from the **web**. It involves **analyzing the data** or **extracting information** about a piece of data from the **web**.

### Types

- Web usage mining  
The process of extracting user information from server logs
- Web structure mining  
The process of using graph theory to analyse the structure of a website
  - Extracting patterns
  - Mining the structure of the document
- Web content mining  
The process of mining, extraction and integration of useful data, information and knowledge from web page content.

## Text Mining Tasks

- text categorization (spam detection, authorship identification, age/gender identification, assigning subject categories, topics, or genres),
- sentiment analysis (detection of attitudes, holder, target, type)
- text clustering (news),
- concept/entity extraction (also called named entity extraction),
- document summarization,
- entity relation modeling (i.e., learning relations between named entities) - useful in question answering.

## Information Retrieval

## Basic Assumptions of Information Retrieval

Information Retrieval (IR) is **finding material** (usually documents) of an **unstructured** nature (usually text) that satisfies an **information need** from within **large collections** (usually stored on computers).

These days we frequently think first of web search, but there are many other cases:

- E-mail search
- Searching your laptop
- Corporate knowledge bases
- Legal information retrieval

**Collection:** A set of documents

Assume it is a static collection for the moment

**Goal:** Retrieve documents with information that is relevant to the user's information need and helps the user complete a task

**Evaluation:** How good are the retrieved docs?

- **Precision:** Fraction of retrieved docs that are relevant to the user's **information need**
- **Recall** : Fraction of relevant docs in collection that are retrieved

## Information Retrieval, an Example

- Which plays of Shakespeare contain the words Brutus AND Caesar but NOT Calpurnia?
- One could grep all of Shakespeare's plays for Brutus and Caesar, then strip out lines containing Calpurnia?
- Why is that not the answer?
  - Slow (for large corpora)
  - NOT Calpurnia is non-trivial
  - Other operations (e.g., find the word Romans near countrymen) not feasible
  - Ranked retrieval (best documents to return)

## Term-document Incidence Matrices

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

So we have a 0/1 vector for each term.

To answer query: take the vectors for Brutus, Caesar and Calpurnia (complemented) → bitwise AND.

110100 AND

110111 AND

101111 =

100100

## Term-document Incidence Matrices for Bigger Collections

- Consider  $N = 1$  million documents, each with about 1000 words.
- Avg 6 bytes/word including spaces/punctuation  
6GB of data in the documents.
- Say there are  $M = 500K$  distinct terms among these.
- Can't build the matrix!
  - $500K \times 1M$  matrix has half-a-trillion 0's and 1's.
  - But it has no more than one billion 1's.  
matrix is extremely sparse.
  - What's a better representation?  
We only record the 1 positions.

## Inverted index

For each term  $t$ , we must store a list of all documents that contain  $t$ . Identify each doc by a **docID**, a document serial number.  
Can we use fixed-size arrays for this?

Brutus	⇒	1	2	4	11	31	45	173	174
Caesar	⇒	1	2	4	5	6	16	57	132
Calpurnia	⇒	2	31	54	101				

What happens if the word Caesar is added to document 14?

## Inverted Index

We need variable-size postings lists

- On disk, a continuous run of postings is normal and best
- In memory, can use linked lists or variable length arrays  
Some trade-offs in size/ease of insertion

Data is organized into:

- Dictionary of terms
- Postings (sorted by DocID)

## Initial stages of text processing

- Tokenization  
Cut character sequence into word tokens
- Normalization  
Map text and query term to same form  
You want U.S.A. and USA to match
- Stemming  
We may wish different forms of a root to match  
authorize, authorization
- Stop words  
We may omit very common words (or not)  
the, a, to, of

## Inverted Index

Query processing: AND

- Locate both words in the Dictionary
- Retrieve their postings
- “Merge” the two postings (intersect the document sets):
  - If postings are **sorted by docID**, it can be done in time  $\mathcal{O}(len1 + len2)$

## Ranked Retrieval

Why do we need more than boolean queries?

- Thus far, our queries have all been Boolean. Documents either match or don't.
- Good for expert users with precise understanding of their needs and the collection.  
Also good for applications: Applications can easily consume 1000s of results.
- Not good for the majority of users.
  - Most users incapable of writing Boolean queries (or they are, but they think it's too much work).
  - Most users don't want to wade through 1000s of results.
  - This is particularly true of web search.
- Boolean queries often result in either too few (0) or too many (1000s) results.
- It takes a lot of skill to come up with a query that produces a manageable number of hits.  
AND gives too few; OR gives too many

## Inverted Index

Query processing: phrase queries

- It no longer suffices to store only term:docs entries
- A first solution: biword indexes
  - Longer phrases can be processed by breaking them down (can lead to false positives)
  - Index blowup due to bigger dictionary
  - Not the standard solution, but can be used as part of a compound strategy
- A second solution: positional indexes
  - In the postings, store, for each term the position(s) in which tokens of it appear
  - For phrase queries, we use a merge algorithm recursively at the document level
  - But we now need to deal with more than just equality

## Ranked Retrieval Models

- Rather than a set of documents satisfying a query expression, in **ranked retrieval models**, the system returns an ordering over the (top) documents in the collection with respect to a query
- **Free text queries**: Rather than a query language of operators and expressions, the user's query is just one or more words in a human language
- When a system produces a ranked result set, large result sets are not an issue

## Scoring as the basis of ranked retrieval

- We wish to return in order the documents most likely to be useful to the searcher
- We need a way of assigning a score to a query/document pair  
Simplified scenario: if the query has only one term
  - If the query term does not occur in the document: score should be 0
  - The more frequent the query term in the document, the higher the score should be

## An Example

Given the following four documents:

- d1:** I bought one computer and one laptop for a great price at one new retail sellers.
- d2:** For the price of this laptop, you could have bought one desktop computer for home and another one for your job.
- d3:** Computer repair budget: 200 euros (a 55 euros discount is offered to registered customers).
- d4:** I need the budget of the laptop and the budget of the printer, in dollars and in euros.

Let us consider a subset of the vocabulary: computer, euros, budget, laptop, price, one, in.

Our search: computer price

## Jaccard coefficient

A commonly used measure of overlap of two sets A and B is the Jaccard coefficient

$$jaccard(A, B) = \frac{A \cap B}{A \cup B}$$

Issues with Jaccard for scoring:

- It doesn't consider term frequency (how many times a term occurs in a document)  
Rare terms in a collection are more informative than frequent terms
- We need a more sophisticated way of normalizing for length

## Binary Term-document Incidence Matrix

	d1	d2	d3	d4
computer	1	1	1	0
euros	0	0	1	1
budget	0	0	1	1
laptop	1	1	0	1
price	1	1	0	0
one	1	1	0	0
in	0	0	0	1

## Term-document Count Matrices

Table: Term frequency  $tf_{t,d}$

	d1	d2	d3	d4
computer	1	1	1	0
euros	0	0	2	1
budget	0	0	1	2
laptop	1	1	0	1
price	1	1	0	0
one	3	1	0	0
in	0	0	0	2

$tf_{t,d}$  = how many times the term  $t$  appears in document  $d$   
Each document is a vector in  $\mathbb{N}^{|V|}$

## Document Frequency

$df_t$  = in how many documents the word  $t$  appears

Table: Document frequency,  $df_t$

Vocabulary	Posting list	$df_t$	$idf_t$
computer	→ d1, d2, d3	3	0.125
euros	→ d3, d4	2	0.301
budget	→ d3, d4	2	0.301
laptop	→ d1, d2, d4	3	0.125
price	→ d1, d2	2	0.301
one	→ d1, d2	2	0.301
in	→ d4	1	0.602

- the base of the log is immaterial
- idf has **no effect** on ranking one term queries.

Inverse document frequency:

$$idf_t = \log_{10}(N/df_t)$$

## Logarithmic Term Frequency

$$w_{t,d} = \begin{cases} 1 + \log_{10} tf_{t,d}, & \text{if } tf_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$

Table: Log-frequency weight,  $w_{t,d}$

	d1	d2	d3	d4
computer	1	1	1	0
euros	0	0	1.301	1
budget	0	0	1	1.301
laptop	1	1	0	1
price	1	1	0	0
one	1.477	1	0	0
in	0	0	0	1.301

Score for a document-query pair: sum over terms  $t$  in  $q$  and  $d$ :

$$\sum_{t \in q \cap d} 1 + \log_{10} tf_{t,d}$$

## Term Frequency, Inverse Document Frequency Weighting

$$tf-idf_{t,d} = (1 + \log_{10} tf_{t,d}) * \log_{10}(N/df_t)$$

Table: Term Frequency, Inverse Document Frequency Weighting:  $tf-idf_{t,d}$

	d1	d2	d3	d4
computer	1 * 0.125	1*0.125	1*0.125	0
euros	0	0	1.301*0.301	1*0.301
budget	0	0	1*0.301	1.301*0.301
laptop	1*0.125	1*0.125	0	1*0.125
price	1*0.301	1*0.301	0	0
one	1.477*0.301	1*0.301	0	0
in	0	0	0	1.301*0.602

Final score:

$$\sum_{t \in q \cap d} (1 + \log_{10} tf_{t,d}) * \log_{10}(N/df_t)$$

Each document is now represented by a real-valued vector of tf-idf weights.

## The Vector Space Model

Documents as vectors

- Now we have a  $|V|$ -dimensional vector space
- Terms are axes of the space
- Documents are points or vectors in this space
- Very high-dimensional: tens of millions of dimensions when you apply this to a web search engine
- These are very sparse vectors - most entries are zero

## Queries as Vectors

- Key idea 1: Do the same for queries: represent them as vectors in the space
- Key idea 2: Rank documents according to their proximity to the query in this space
- proximity = similarity of vectors
- Recall: We do this because we want to get away from the ~~you're-either-in-or-out~~ Boolean model
- Instead: rank more relevant documents higher than less relevant documents

## Formalizing Vector Space Proximity

- Using Euclidean distance is a bad idea, because Euclidean distance is large for vectors of different lengths
- Using angle instead of distance is much better:
  - take a document  $d$  and append it to itself. Call this document  $d'$ .
  - "Semantically"  $d$  and  $d'$  have the same content
  - The Euclidean distance between the two documents can be quite large
  - The angle between the two documents is 0, corresponding to maximal similarity.

## The Cosine Similarity

### From angles to cosine

- The following two notions are equivalent.
  - Rank documents in decreasing order of the angle between query and document
  - Rank documents in increasing order of cosine(query,document)
- Cosine is a monotonically decreasing function in  $[0^\circ, 180^\circ]$

### Length normalization

- A vector can be (length-) normalized by dividing each of its components by its length - for this we use the  $L_2$  norm:  
$$\|\vec{x}\|_2 = \sqrt{\sum_i x_i^2}$$
- Dividing a vector by its  $L_2$  norm makes it a unit (length) vector (on surface of unit hypersphere)
- Effect on the two documents  $d$  and  $d'$  ( $d$  appended to itself): they have identical vectors after length-normalization.  
Long and short documents now have comparable weights.

## Cosine Similarity

$$\cos(q, d) = \frac{\sum_{t \in q \cap d} q_t d_t}{\sqrt{\sum_{t \in V} q_t^2} * \sqrt{\sum_{t \in V} d_t^2}}$$

$q_t$  is the tf-idf weight of term  $t$  in the query

$d_t$  is the tf-idf weight of term  $t$  in the document

## Variants of tf-idf

Term frequency	Document frequency	Normalization
n (natural): $tf_{t,d}$	n (no): 1	n (none): 1
l (logarithm): $1 + \log_{10}(tf_{t,d})$	t (idf): $\log(\frac{N}{df_t})$	c (cosine): $\frac{1}{\sqrt{\sum_i x_i^2}}$
a (augmented): $0.5 + \frac{0.5 * tf_{t,d}}{\max_t(tf_{t,d})}$	p (prob idf): $\max\{0, \log \frac{N - df_t}{df_t}\}$	u (pivoted unique): $1/u$
b (boolean) $\begin{cases} 1 & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$		b (byte size): $1/CharLength^\alpha$
L (log ave): $\frac{1 + \log_{10}(tf_{t,d})}{1 + \log_{10}(\text{ave}_{t \in d}(tf_{t,d}))}$		( $\alpha < 1$ )

## Weighting Schemes

- Weighting may differ in queries vs documents
- Many search engines allow for different weightings for queries vs. documents
- SMART Notation: denotes the combination in use in an engine, with the notation ddd.qqq, using the acronyms from the previous table
- A very standard weighting scheme is: Inc.ltc

$$\cos(q, d) = \frac{\sum_{t \in q \cap d} q_t d_t}{\sqrt{\sum_{t \in V} q_t^2} * \sqrt{\sum_{t \in V} d_t^2}}$$

$$q_t = (1 + \log_{10} tf_{t,q}) * \log_{10}(N/df_t)$$

$$d_t = (1 + \log_{10} tf_{t,d}) * \log_{10}(N/df_t)$$

## Evaluating an IR System

- An information need is translated into a query
- Relevance is assessed relative to the information need not the query
- E.g., Information need: *I'm looking for information on whether drinking red wine is more effective at reducing your risk of heart attacks than white wine.*
- Query: *wine red white heart attack effective*
- You evaluate whether the doc addresses the information need, not whether it has these words



## Two Current Evaluation Measure

### Precision-recall curve

(Assume 10 relevant documents in collection)

		Recall	Precision
1	R	0.1	1.0
2	N	0.1	0.5
3	N	0.1	0.33
4	R	0.2	0.5
5	R	0.3	0.6
6	N	0.3	0.5
7	R	0.4	0.57
8	N	0.4	0.5

### Mean average precision (MAP)

Average of the precision value obtained for the top  $k$  documents, each time a relevant document is retrieved.

$$MAP = (1/1 + 2/4 + 3/5 + 4/7)/4 = 0.67$$